

# Predictive Modelling - V

## Ensemble Models

---

Rita P. Ribeiro

Data Mining I - 2023/2024



## Summary

- Ensembles: Motivation
- Types Ensembles
- Ensemble Methods
  - Random Forest
  - AdaBoost
  - XGBoost

# Ensembles

---

## Predictive Modelling: Where we at?

- Probabilistic Approaches
  - e.g. Naive Bayes, Bayesian Networks
- Mathematical Formulae
  - e.g. multiple linear regression
- Logical Approaches
  - e.g. CART
- Distance-based Approaches
  - e.g. kNN
- Optimization Approaches
  - e.g. SVM, ANN
- **Ensemble Approaches**
  - e.g. Random Forest, XgBoost

## Ensemble Models

- **Ensembles**: collections of models that are used together to address a certain prediction problem
- Different learning algorithms exploit:
  - different languages for representing generalizations of the examples;
  - different search spaces;
  - different evaluation functions of the hypothesis;
- For complex problems it is hard to find a model that “explains” all observed data
- There is no overall better algorithm → *No free lunch theorem*

## Ensemble Models

- Averaging over a set of models typically leads to significantly better results, given certain conditions.
- An ensemble of classifiers improves over individual classifiers iif (Dietterich 2002):
  - they perform better than random guess;
  - they have non-correlated errors;
  - they commit errors in different regions of the instance space.


How to achieve such **diversity**?

- Combining outputs in different ways
- Perturbing the set of training examples
  - Homogeneous Models (Bagging, Boosting)
  - Heterogeneous Models (Cascading, Stacking)
- Perturbing the set of attributes

## Bias-Variance Trade-off

### The Bias-Variance Decomposition of Prediction Error


- The prediction error of a model can be split in two main components: the bias and the variance components
- **bias**: error that is due to the poor ability of the model to fit the seen data
- **variance**: error related to the sensibility of the model to the given training data



## Bias-Variance Trade-off

When learning a prediction model, there is **bias-variance trade-off**.

- Decreasing the bias by adjusting more to the training sample → higher variance - the over-fitting phenomenon
- Decreasing the variance by being less sensitive to the given training data → higher bias



## Bias-Variance Trade-off


- Ensembles are able to reduce both components of the error
- Their approach consists on:
  - applying the same algorithm to different samples of the data
  - use the resulting models in a voting/averaging schema to obtain predictions for new cases

# Types of Ensembles

---

## Types of Ensembles

- Independent or Parallel Models
- Coordinated or Sequential Models



## Types of Ensembles: Independent or Parallel Models

- Construct the models independently in a way that ensures some diversity among them
- How to reach diversity?
  - applying the models on somewhat different training sets
  - applying the models on data sets using different predictors

## Types of Ensembles: Coordinated or Sequential Models

- Construct a “larger” model by composing it from smaller and integrated models
- Each individual model has a weighted participation in the ensemble predictions
- What is the right component models and their respective weight to achieve a good predictive performance?

# Ensemble Methods

---

## Ensembles using Independent Models: Bagging

**Bagging** or Bootstrap Aggregating (Breiman 1996)

- Method that obtains a set of  $k$  models using different bootstrap samples of the given training data
  - sample with replacement of the same size as the available data
  - for each learner, a small proportion of examples will be different
- If the base learner has a high variance (i.e. very sensitive to variations on the training sample), this will ensure diversity among the  $k$  models
- Bagging should be applied to base learners with high variance



## Ensembles using Independent Models: Bagging

- Requires unstable algorithms (greedy like)
- Algorithms sensible to small perturbations of the training set;
  - Decision trees, Rule learners, etc.
- Easy to implement with any algorithm.
- Easy to implement in parallel environments.
- The bias-variance argument:
  - Error decreases due to reduction in the variance component.

## Ensembles using Independent Models: Random Forests

### Varying the Predictors

- Another way of generating a diverse set of models is by using different randomly chosen predictors
- The idea is similar to bagging but instead of generating samples of the cases we generate samples of the variables

### Random Forests (Breiman 2001)

- Combine the ideas of bagging together with the idea of random selection of predictors
- Set of tree-based models where each tree is obtained from a bootstrap sample of the original data and uses random selection of variables during tree growth

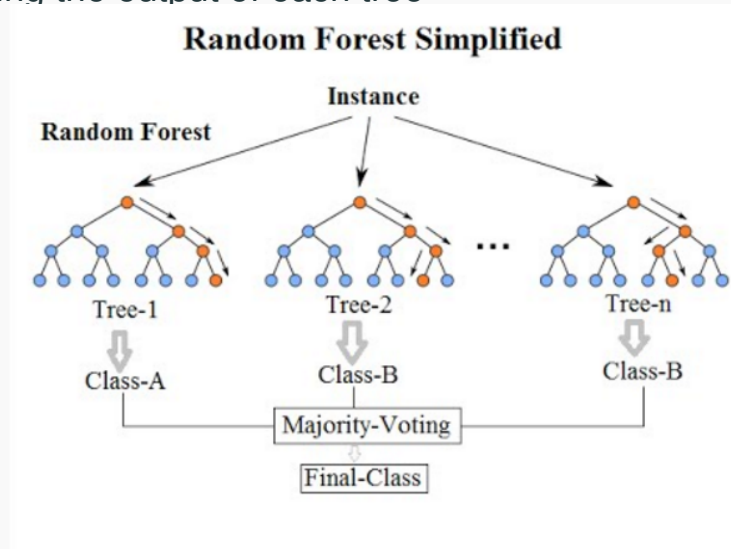
### Learning phase (main idea)

- For  $t = 1$  to  $T$ ,  $T$  is number of trees
  - draw a random sample with replacement from the training set  $D_t$
  - train a tree model  $h_t(\mathbf{x})$  on  $D_t$  without pruning
  - at each candidate split in the learning process, uses a random subset of the  $m$  features.
- Return  $\{h_t(\mathbf{x}) | 1 \leq t \leq T\}$

## Ensembles using Independent Models: Random Forests

### Prediction phase

- Predict the class obtained by majority vote, or the value by averaging the output of each tree




## Ensembles using Independent Models: Random Forests

### Other Uses of Random Forests: Variable Importance

- Which variables have the most predictive power?
- Two importance measures:
  - how much the accuracy decreases / mean square error increases when the variable is excluded
  - how much the impurity decreases when the variable is chosen to split a node.

## Ensembles using Independent Models: Random Forests

### Variable Importance: an example



## Ensembles using using Independent Models: Random Forests

### Hyperparameters

- Number of trees
  - recommended number of trees is 1000.
  - to obtain more reliable statistics for the attribute importance, 5000 trees are recommended.
- Number of attributes to randomly select at each node
  - it must be tuned
  - its optimum value is problem dependent.
  - rule of thumb:  $\sqrt{p}$ ,  $p$  is the number of predictive attributes

## Ensembles using Independent Models: Random Forests

### Pros

- Do not require elaborate tuning of the hyper-parameters. Often these can/should be optimized.
- The most important parameter to tune is the number of trees to grow, typically the larger the best.
- Do not need to worry about creating very complex trees.

### Cons

- Do not provide the interpretability level of a Decision Tree

## Ensembles using Coordinated Models: Boosting

### Boosting (Schapire 1990)

Can a set of weak learners create a single strong learner?

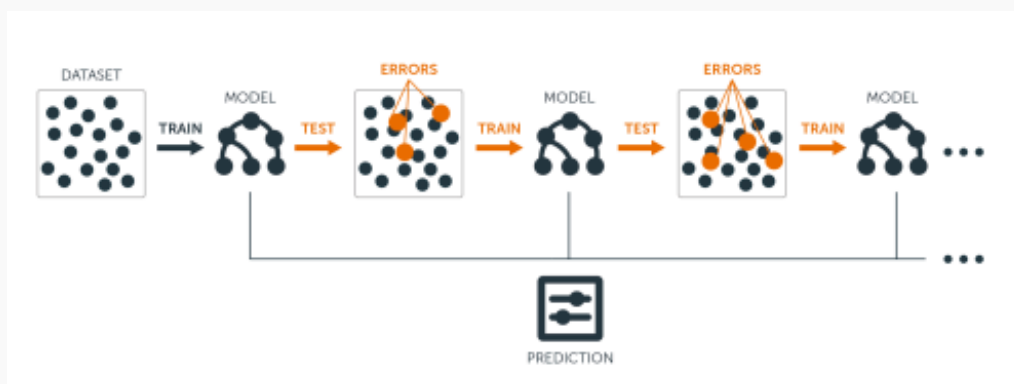
- A “weak” learner is a model that alone is unable to correctly approximate the unknown predictive function
- A “strong” learner has that ability
- Boosting algorithms work by iteratively creating a strong learner by adding at each iteration a new weak learner to make the ensemble
- Weak learners are added with weights that reflect the learner’s predictive power

## Ensembles using Coordinated Models: Boosting

- After each addition the data is re-weighted such that cases that are still poorly predicted gain more weight
- The weight indicates the probability of the example being select in a uniform sampling;
- This means that each new weak learner will focus on the errors of the previous ones
- It fits many real-world problems, where observed examples tend to have different learning difficulty levels.
  - e.g. examples close to the decision surface are typically more difficult

## Ensembles using Coordinated Models: Boosting

- The prediction: weighted voting/average of each learner.



## Ensembles using Coordinated Models: Boosting

Three ways through which boosting can be carried out:


- Adaptive Boosting or AdaBoost (Freund and Schapire 1996)
- Gradient Boosting Machine (Friedman 2000)
- XGBoost (Chen and Guestrin 2016)

## Ensembles using Coordinated Models: AdaBoost

**AdaBoost** or Adaptive Boosting (Freund and Schapire 1996)

- **Iterative** process: new models are added to form an ensemble
- **Adaptive**: at each new iteration of the algorithm, the new models are built to try to overcome the errors made in the previous iterations
- At each iteration the weights of the training cases are adjusted
- Cases wrongly predicted get their weight increased to make new models focus on accurately predicting them
- The main hyperparameter is **number of iterations**
- AdaBoost was created for classification although variants for regression exist

## Ensembles using Coordinated Models: AdaBoost



Source: <https://medium.com/divyagera2402>

## Ensembles using Coordinated Models: Boosting

### The Algorithm (main idea)


- Start with uniform weights  $w_i^{(1)} = 1/|D|$  for all  $\mathbf{x}_i \in D$
- For  $t = 1$  to  $T$ 
  - build weak model  $h_t(\mathbf{x})$
  - calculate weighted error  $e_t = \sum_i w_i^{(t)} I(y_i \neq h_t(\mathbf{x}_i))$
  - the weight of this weak model:  $\alpha_t = \frac{1}{2} \ln \left( \frac{1-e_t}{e_t} \right)$
  - update case weights  $w_i^{(t+1)} = \frac{w_i^{(t)} \exp(-\alpha_t I(y_i \neq h_t(\mathbf{x}_i)))}{Z_t}$  where  $Z_t$  is chosen to make all  $w_i^{(t+1)}$  sum up to 1
- Return a form of additive model composed of  $t$  weak models

$$H(\mathbf{x}) = \sum_{t=1}^T \alpha_t h_t(\mathbf{x})$$



## Ensembles using Coordinated Models: Boosting

Evolution of the error as you increase the number of weak learners.



## Ensembles using Coordinated Models: GBM

### Gradient Boosting Machine (Friedman 2000)

- Sequential ensemble learning
- Contrary to AdaBoost, it does not adjust the examples weights at every iteration
- It fits the new learner to the residual errors made by the previous learner
- The present learner is always more effective than the previous one
- Goal: at each step, adds a weak learner to increase the performance and build a strong learner.

## Ensembles using Coordinated Models: GBM

- Re-defines boosting as a **numerical optimization problem**
- Objective: minimize the loss function of the model by adding weak learners using a gradient-descent procedure.
- Major difference: how it identifies the shortcomings of weak learners (e.g. decision trees).
- It uses gradients in the loss function as a measure indicating how good are model's coefficients are at fitting the underlying data
- Like AdaBoost, it can be used for both classification and regression problems

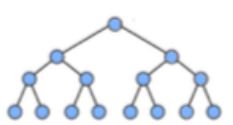
## Ensembles using Coordinated Models: GBM

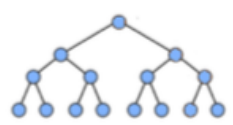
### Learning phase (main idea)

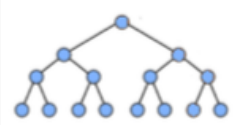
- Build an additive tree model by adding new trees to complement the already-built ones

#	$\mathbf{x}$	$y$	#	$\mathbf{x}$	$y$	#	$\mathbf{x}$	$y$
1	$\mathbf{x}_1$	$y_1$	1	$\mathbf{x}_1$	$\ell(y_1, h_1(\mathbf{x}_1))$	1	$\mathbf{x}_1$	$\ell(y_1, h_1(\mathbf{x}_1) + h_2(\mathbf{x}_1))$
2	$\mathbf{x}_2$	$y_2$	2	$\mathbf{x}_2$	$\ell(y_2, h_1(\mathbf{x}_2))$	2	$\mathbf{x}_2$	$\ell(y_2, h_1(\mathbf{x}_2) + h_2(\mathbf{x}_2))$
3	$\mathbf{x}_3$	$y_3$	3	$\mathbf{x}_3$	$\ell(y_3, h_1(\mathbf{x}_3))$	3	$\mathbf{x}_3$	$\ell(y_3, h_1(\mathbf{x}_3) + h_2(\mathbf{x}_3))$
...	...	...	...	...	...	...	...	...

$h_1$ : 

$h_2$ : 

$h_3$ : 

## Ensembles using Coordinated Models: GBM

Learning phase (main idea) - cont.


• Objective: minimize  $Obj = \sum_{i=1}^n \ell(y_i, \hat{y}_i) + \sum_{k=1}^K \mathcal{R}(h_k)$   
where

- $\sum_{i=1}^n \ell(y_i, \hat{y}_i)$  is the training loss
  - measures how well the model fits on training data
- $\sum_{k=1}^K \mathcal{R}(h_k)$  is the regularization term
  - measures the complexity of trees (nr of leafs and  $L_2$ -norm of leaf scores)

## Ensembles using Coordinated Models: GBM

Prediction phase

- Response is the optimal linear combination of all decision trees



## Ensembles using Coordinated Models: GBM

### Hyperparameters

- **Learning rate** ( $\alpha$ ) is a multiplying factor on the errors for the subsequent trees.
  - It controls how fast the model learns: the lower  $\alpha$ , the slower the model learns.
  - The advantage of slower learning rate: the model becomes more robust and avoids overfitting.
  - However, learning slowly comes at a cost: it takes more time to train the model
- **Number of trees** used in the model.
  - If the learning rate is low, we need more trees to train the model.
  - However, it creates a high risk of overfitting to use too many trees.

## Ensembles using Coordinated Models: XGBoost


### eXtreme Gradient Boosting (XGBoost) (Chen and Guestrin 2016)

- An advanced version of Gradient Boosting Method
- Software and hardware optimization
  - a scalable tree boosting system
- Some features:
  - clever penalisation of trees: weights of the trees that are calculated with less evidence is shrunk more heavily
  - extra randomisation parameter to reduce correlation between trees
  - parallelization, cache optimization, distributed computing, etc.

## Ensembles using Coordinated Models: XGBoost

### Feature Importance

- How useful each feature was in the construction of the boosted decision trees?
- The more the feature is selected for splitting, the higher its relative importance.
- Importance is calculated for a single decision tree by number of times the feature is selected for splitting, weighted by the improvement to the model as a result of each split.
- The feature importances are then averaged across all of the the decision trees within the model.



## Ensemble Methods: Wrap-up

- Well designed ensembles of predictive models allow improvement of performance over their individual elements.
- Necessary conditions:
  - variability between elements;
  - low error correlation;
  - each individual model must be better than a random choice

### Bagging Methods

- Error reduction due to reduction in variance;
- Effective with unstable models;

### Boosting Methods

- Error reduction due to reduction in bias and variance;
- Risky in problems with noise (increase of the error);

## References

---

- Breiman, Leo. 1996. "Bagging Predictors." *Machine Learning* 24 (2): 123–40.
- . 2001. "Random Forests." *Machine Learning* 45 (1): 5–32.
- Chen, Tianqi, and Carlos Guestrin. 2016. "XGBoost: A Scalable Tree Boosting System." In *22nd ACM SIGKDD*, 785–94.
- Dietterich, Thomas. 2002. "Ensemble Learning." In *The Handbook of Brain Theory and Neural Networks*, 2nd ed., 405–8. MIT Press.
- Flach, Peter. 2012. *Machine Learning: The Art and Science of Algorithms That Make Sense of Data*. Cambridge University Press.
- Freund, Yoav, and Robert E. Schapire. 1996. "Experiments with a New Boosting Algorithm." In *13th ICML*, 148–56. Morgan Kaufmann.
- Friedman, Jerome H. 2000. "Greedy Function Approximation: A Gradient Boosting Machine." *Annals of Statistics* 29: 1189–1232.
- Gama, João. 2018. "KDD Course." Slides.
- Schapire, Robert E. 1990. "The Strength of Weak Learnability." *Machine Learning* 5 (2): 197–227.
- Torgo, Luís. 2017. "Data Mining i Course." Slides.