

Computer Vision – TP12

Advanced Segmentation

Miguel Coimbra, Hélder Oliveira

Outline

- Segmentation by Fitting
- Active Contours
- Semantic Segmentation

Topic: Segmentation by Fitting

- Segmentation by Fitting
- Active Contours
- Semantic Segmentation

Fitting and Clustering

- Another definition for segmentation:
 - Pixels belong together because they conform to some model
- Sounds like “Segmentation by Clustering”...
- Key difference:
 - The model is now **explicit**

We have a mathematical model for the object we want to segment.
Ex: A line

Hough Transform

- Elegant method for direct object recognition
- Edges need not be connected
- Complete object need not be visible
- Key Idea: Edges **VOTE** for the possible model

Image and Parameter Spaces

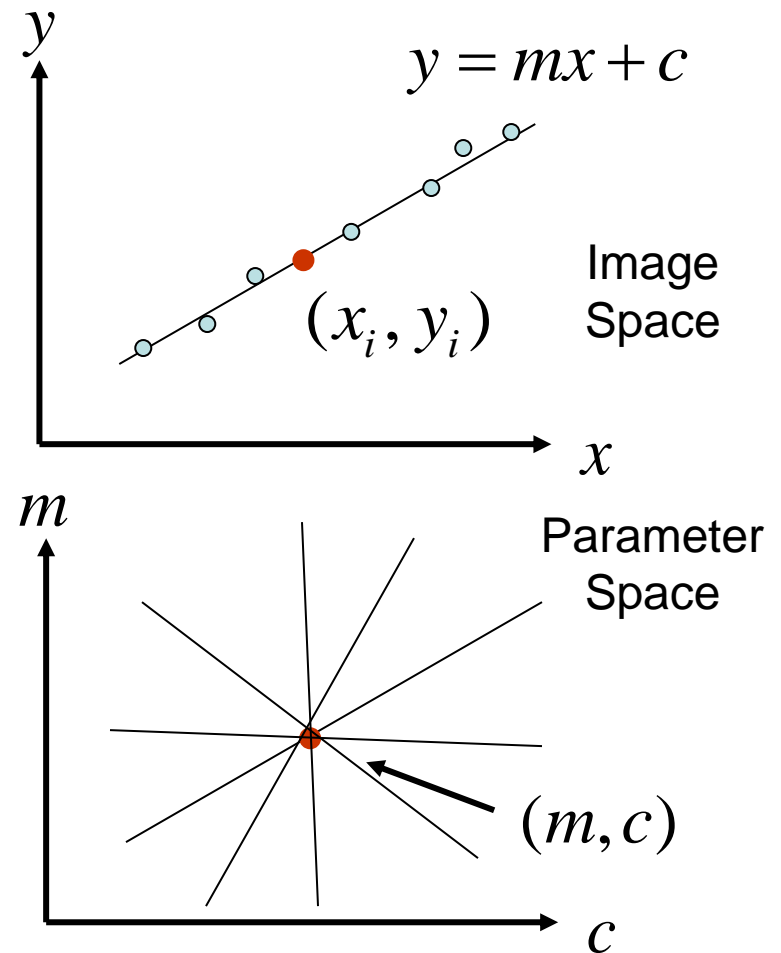
Equation of Line: $y = mx + c$

Find: (m, c)



Consider point: (x_i, y_i)

$$y_i = mx_i + c \quad \text{or} \quad c = -x_i m + y_i$$



Parameter space also called Hough Space



Example - Lines



Example - Circles





Least Squares Line Fitting

- Popular fitting procedure
- Simple but biased (why?)
- Consider a line:

$$y = ax + b$$

- What is the line that best predicts all observations (x_i, y_i) ?
 - Minimize: $\sum_i (y_i - ax_i - b)^2$





This looks better...

Total Least Squares

- Works with the actual distance between the point and the line (rather than the vertical distance)
- Lines are represented as a collection of points where:

$$ax + by + c = 0$$

- And:

$$a^2 + b^2 = 1$$

Again... Minimize the error, obtain the line with the 'best fit'.

Point correspondence

- We can estimate a line but, **which points are on which line?**
- Usually:
 - We are fitting lines to edge points, so...
 - Edge directions can give us hints!
- What if I only have isolated points?
- Let's look at two options:
 - Incremental fitting
 - Allocating points to lines with K-means

Incremental Fitting

- Start with connected *curves* of edge points
- Fit *lines* to those points in that curve
- Incremental fitting:
 - Start at one end of the *curve*
 - Keep fitting all points in that curve to a line
 - Begin another line when the fitting deteriorates too much
- Great for closed curves!

```
Put all points on curve list, in order along the curve
empty the line point list
empty the line list
```

```
Until there are two few points on the curve
  Transfer first few points on the curve to the line point list
  fit line to line point list
```

```
  while fitted line is good enough
    transfer the next point on the curve
    to the line point list and refit the line
  end
```

```
  transfer last point back to curve
  attach line to line list
end
```

K-means allocation

- What if points carry no hints about which line they lie on?
- Assume there are k lines for the x points.
- Minimize:
$$\sum_{\text{lines}} \sum_{\text{points}} \text{dist}(\text{line}, \text{point})^2$$
- Iteration:
 - Allocate each point to the closest line
 - Fit the best line to the points allocated to each line

Hypothesize k lines (perhaps uniformly at random)
or

hypothesize an assignment of lines to points
and then fit lines using this assignment

Until convergence


- allocate each point to the closest line
- refit lines

Topic: Active Contours

- Segmentation by Fitting
- **Active Contours**
- Semantic Segmentation


Active Contours

- Given: initial contour (model) near desired object



Active Contours



- Goal: evolve the contour to fit exact object boundary






- How?
 - Reward solutions next to high image gradients
 - Punish solutions that deform shape too much
 - Iteratively find the ‘best’ solution to these requirements

Intuition - Elastic Band

- Contour evolves to a low-energy solution, but is hindered by obstacles
- Better intuition: Gravity
 - Contour is ‘attracted’ to specific image features
 - Contour resists to any deformation of its shape



Strong motivation – Moving deformable objects




Things we need to consider

- Representation of the contours
- Defining the energy functions
 - External
 - Internal
- Minimizing the energy function

Representation


- We'll consider a discrete representation of the contour, consisting of a list of 2d point positions ("vertices")



$$v_i = (x_i, y_i),$$

for $i = 0, 1, \dots, n-1$


- At each iteration, we'll have the option to move each vertex to another nearby location ("state")



Energy function

The total energy (cost) of the current snake is defined as:


$$E_{total} = E_{external} + E_{internal}$$



- External energy: encourage contour to fit on places where specific image structures exist
- Internal energy: encourage prior shape preferences


A good fit between the current deformable contour and the target shape in the image will yield a low value for this cost function

External image energy



Magnitude of gradient

$$G_x(I)^2 + G_y(I)^2$$




- (Magnitude of gradient)

$$-\left(G_x(I)^2 + G_y(I)^2\right)$$

External image energy

- Gradient images $G_x(x, y)$ and $G_y(x, y)$



- External energy at a point on the curve is:


$$E_{external}(v) = -(|G_x(v)|^2 + |G_y(v)|^2)$$

- External energy for the whole curve:

$$E_{external} = - \sum_{i=0}^{n-1} |G_x(x_i, y_i)|^2 + |G_y(x_i, y_i)|^2$$

Internal energy

For a *continuous* curve, a common internal energy term is the “bending energy”



At some point $v(s)$ on the curve, this is:


$$E_{internal}(v(s)) = \alpha \left| \frac{dv}{ds} \right|^2 + \beta \left| \frac{d^2v}{ds^2} \right|^2$$

Tension,
Elasticity Stiffness,
Curvature

Internal energy

- For our discrete representation:

$$\mathbf{v}_i = (x_i, y_i) \quad i = 0 \dots n-1$$



$$\frac{d\mathbf{v}}{ds} \approx \mathbf{v}_{i+1} - \mathbf{v}_i \quad \frac{d^2\mathbf{v}}{ds^2} \approx (\mathbf{v}_{i+1} - \mathbf{v}_i) - (\mathbf{v}_i - \mathbf{v}_{i-1}) = \mathbf{v}_{i+1} - 2\mathbf{v}_i + \mathbf{v}_{i-1}$$

- Internal energy for the whole curve:

$$E_{internal} = \sum_{i=0}^{n-1} \alpha \|\mathbf{v}_{i+1} - \mathbf{v}_i\|^2 + \beta \|\mathbf{v}_{i+1} - 2\mathbf{v}_i + \mathbf{v}_{i-1}\|^2$$


Example: compare curvature

$$\begin{aligned} E_{curvature}(v_i) &= \|v_{i+1} - 2v_i + v_{i-1}\|^2 \\ &= (x_{i+1} - 2x_i + x_{i-1})^2 + (y_{i+1} - 2y_i + y_{i-1})^2 \end{aligned}$$

● (2,5)



$$\begin{aligned} (3 - 2(2) + 1)^2 + (1 - 2(5) + 1)^2 \\ = (-8)^2 = 64 \end{aligned}$$



$$\begin{aligned} (3 - 2(2) + 1)^2 + (1 - 2(2) + 1)^2 \\ = (-2)^2 = 4 \end{aligned}$$

Penalizing elasticity

- Current elastic energy definition uses a discrete estimate of the derivative:

$$E_{elastic} = \sum_{i=0}^{n-1} \alpha \|v_{i+1} - v_i\|^2$$

- This rewards very small shapes!
- Instead -> Reward an 'average distance \bar{d} between pairs of points'

$$= \alpha \cdot \sum_{i=0}^{n-1} \left((x_{i+1} - x_i)^2 + (y_{i+1} - y_i)^2 - \bar{d} \right)^2$$

Total energy


$$E_{total} = E_{internal} + \gamma E_{external}$$

$$E_{external} = - \sum_{i=0}^{n-1} |G_x(x_i, y_i)|^2 + |G_y(x_i, y_i)|^2$$


$$E_{internal} = \sum_{i=0}^{n-1} \alpha \left(\bar{d} - \|v_{i+1} - v_i\| \right)^2 + \beta \|v_{i+1} - 2v_i + v_{i-1}\|^2$$

Energy weights


- e.g., α weight controls the penalty for internal elasticity



large α




medium α



small α


Energy minimization: greedy

- For each point, search window around it and move to where energy function is minimal
 - Typical window size, e.g., 5 x 5 pixels
- Stop when predefined number of points have not changed in last iteration, or after max number of iterations
- Note:
 - Convergence not guaranteed
 - Need decent initialization



Tracking via deformable contours

1. Use final contour/model extracted at frame t as an initial solution for frame $t+1$
2. Evolve initial contour to fit exact object boundary at frame $t+1$
3. Repeat, initializing with most recent frame



Deformable contours: pros and cons

Pros:

- Useful to track and fit non-rigid shapes
- Contour remains connected
- Possible to fill in “subjective” contours
- Flexibility in how energy function is defined, weighted.

Cons:

- Must have decent initialization near true boundary, may get stuck in local minimum
- Parameters of energy function must be set well based on prior information


Topic: Semantic Segmentation

- Segmentation by Fitting
- Active Contours
- **Semantic Segmentation**

Remember 'Segmentation'?

- Separation of the image in different areas
 - Objects
 - **Areas with similar visual or semantic characteristics**


First form regions based on visual characteristics, then find the semantics of each region



Semantic Segmentation

- Separation of the image in different areas
 - Objects
 - **Areas with similar visual or semantic characteristics**

First classify each pixel, and only then form regions (much harder!!)



Classification and Segmentation

Classification




CAT

**Classification
+ Localization**



CAT

Object Detection



CAT, DOG, DUCK

**Instance
Segmentation**



CAT, DOG, DUCK

Single object

Multiple objects

Source <http://cs224d.stanford.edu/index.html>

Semantic Segmentation

Other Computer Vision Tasks

Semantic Segmentation



GRASS, CAT,
TREE, SKY

No objects, just pixels

**Classification
+ Localization**



CAT

Single Object

**Object
Detection**



DOG, DOG, CAT

Multiple Object

**Instance
Segmentation**



DOG, DOG, CAT

This image is CC0 public domain

Fei-Fei Li & Justin Johnson & Serena Yeung

Lecture 11 - 17 May 10, 2017

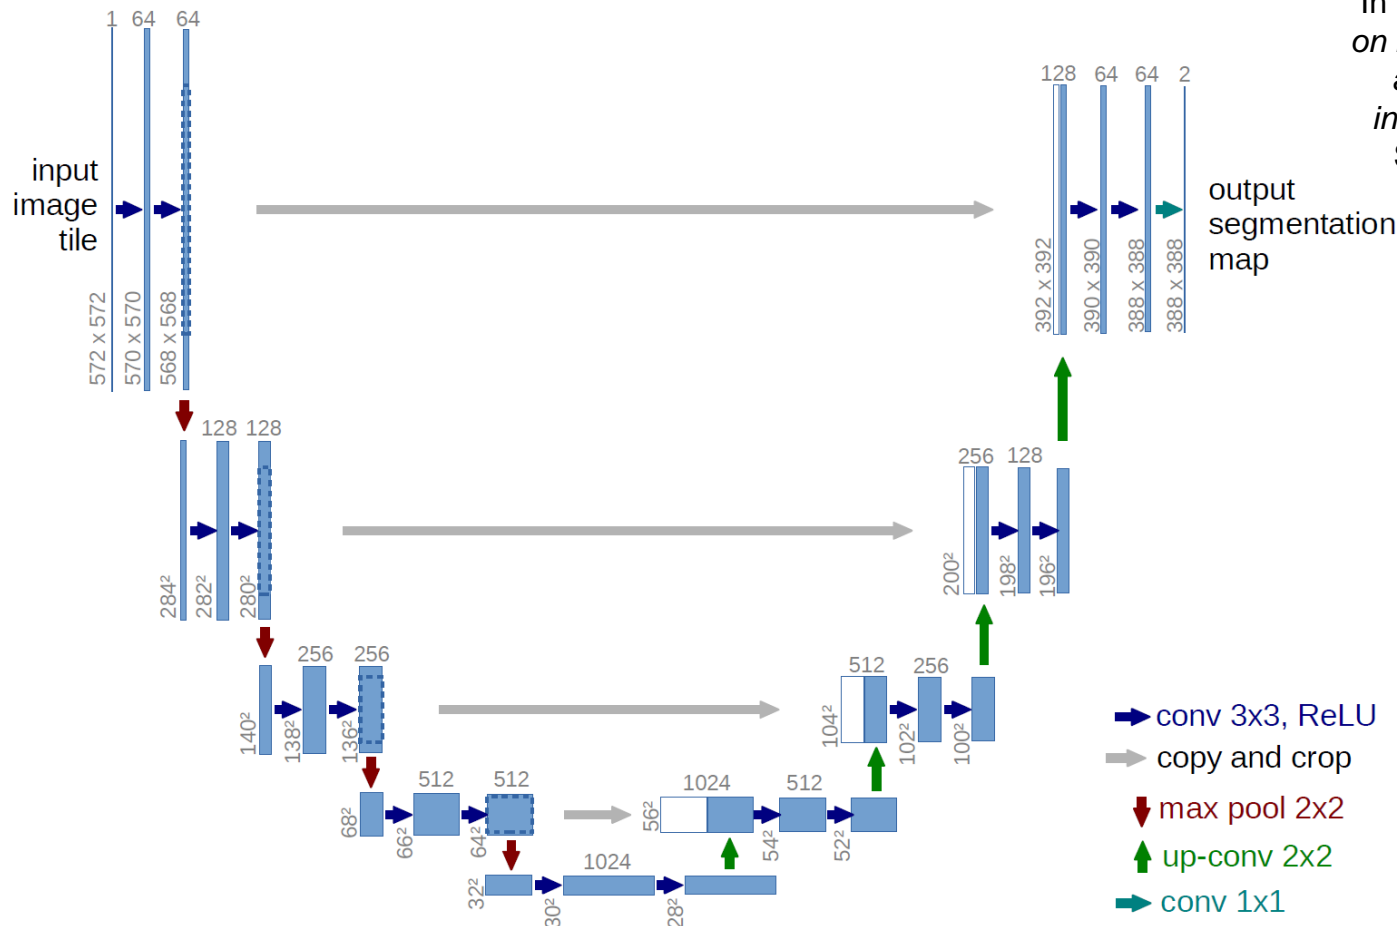
Source http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture11.pdf

Semantic Segmentation

- Requires sophisticated pixel-level classification algorithms to be effective
- Powerful data-based approach to segmentation
- Fueled by recent advances in deep neural networks, such as U-NET

U-Net

O. Ronneberger, P. Fischer, and T. Brox. "U-net: Convolutional networks for biomedical image segmentation." In *International Conference on Medical image computing and computer-assisted intervention*, pp. 234-241. Springer, Cham, 2015.



- Encoder-decoder structure

Resources

- Szeliski, “Computer Vision: Algorithms and Applications”, Springer, 2022
 - Chapter 6 – “Recognition”
 - Chapter 7 – “Feature Detection and Matching”