

# Big Data and Cloud Computing, 23/24 (Part 2)

Inês Dutra and Eduardo Marques  
DCC-FCUP  
ines@dcc.fc.up.pt (room: 1.31)

April 10th, 2024

# Data Mining and Machine Learning: recap

- Learning?

- ▶ “An agent **learns** if it improves its performance in future tasks after making observations about the past or current world.” (Tom Mitchel)

# Machine Learning: very brief overview

- Learning?

- ▶ Given observations  $O$ ,  
described by features  $f_1, f_2, \dots, f_n$ ,  
the task of a machine learning algorithm is:
  - to find patterns based on features  $f_1, f_2, \dots, f_n$  (all or some of them),  
that distinguish among different groups of observations OR
  - to find a function that will **predict** new observations

# Machine Learning: very brief overview

- Learning?

- ▶ Can be **supervised**:

- Given features  $f_1, f_2, \dots, f_n$ ,  
**and** a special feature, the **target** variable (ground truth),  
find a model that can **predict** the target variable for **new** observations  
that are described by features  $f_1, f_2, \dots, f_n$
    - The supervised learning task can be **classification** or **regression**

- ▶ Can be **unsupervised**: find subgroups of patterns, no target variable is known or provided

- clustering
    - association rules

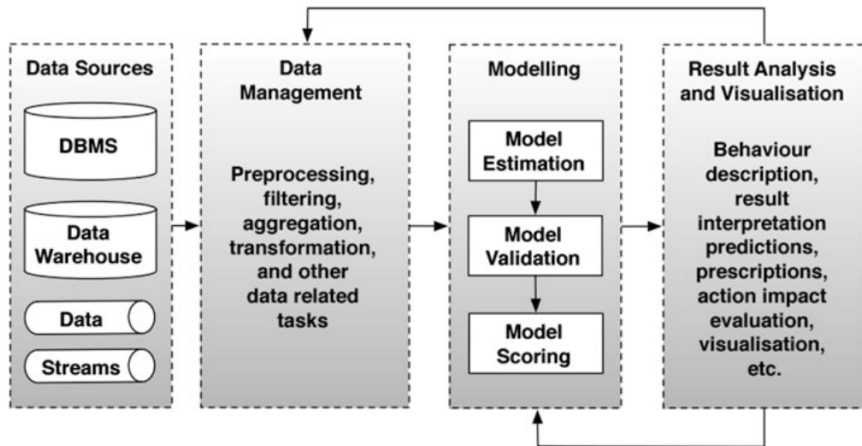
- ▶ Other learning methods: reinforcement learning, matrix factorization for recommender systems

- ▶ *background/prior knowledge*: description of observations, necessary to improve the learning

# Data Mining and Machine Learning: recap

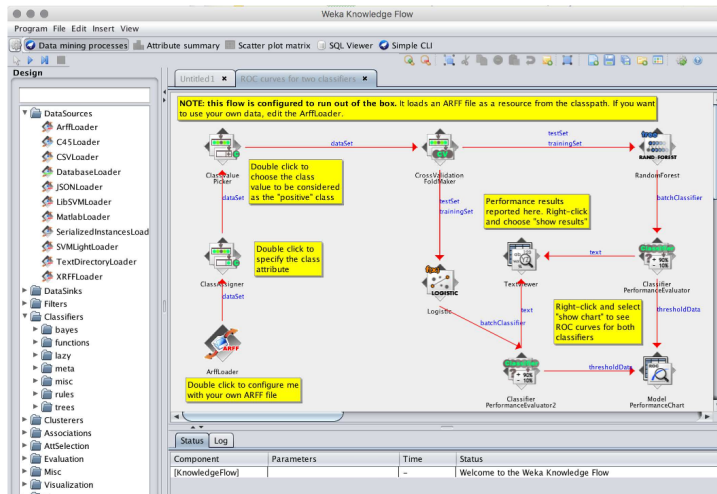
- Workflow (Dataflow - Knowledgeflow):
  - ▶ Data preprocessing
    - transformation: normalization, standardization, averaging, median, denoising, filtering
    - preparation: depends on the task, algorithm, package or library being used
  - ▶ Machine learning task, algorithm
  - ▶ Validation: cross-validation, bootstrapping
- Workflow tools: WEKA KnowledgeFlow, RapidMiner, Orange3, Taverna, Condor DAGMan, Pegasus, Google Dataflow, Google Composer (Apache Airflow)

# Data Mining and Machine Learning: workflow



Assunção, M.D., Calheiros, R.N., Bianchi, S., Netto, M.A.S., Buyya, R.: Big data computing and clouds: trends and future directions. J. Parallel Distrib. Comput. 79–80, 3–15 (2015)

# Example of workflow in WEKA



`java -jar weka.jar ⇒ KnowledgeFlow`

# Example of workflow with Orange3

(installation needed, go to <https://orange.biolab.si/>)




# Limitations

- Most systems and tools for data analysis are not scalable
- I/O, memory, computing power


- Computational resources: memory, CPU, I/O, storage
- I/O:
  - ▶ **Experiment 1:** SPSS, MatLab, R and DataIP (in-house implementation)
    - dataset of patients, originally 200K entries, 6 numeric variables without nulls
    - varying sizes: 300, 200k, 1M, 5M, 10M
  - ▶ **Experiment 2:** job that needs to fetch data files from a remote site

# Scalability: Experiment 1, I/O



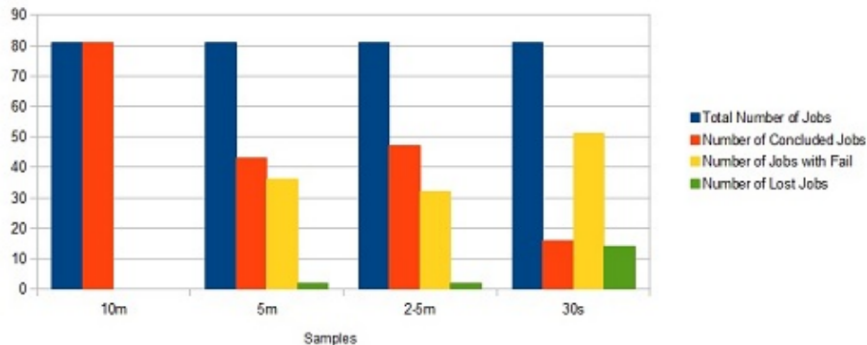
# Scalability: Experiment 1, simple computing: summary

Summary



## Scalability: Experiment 2, file transfer

Jobs x Samples



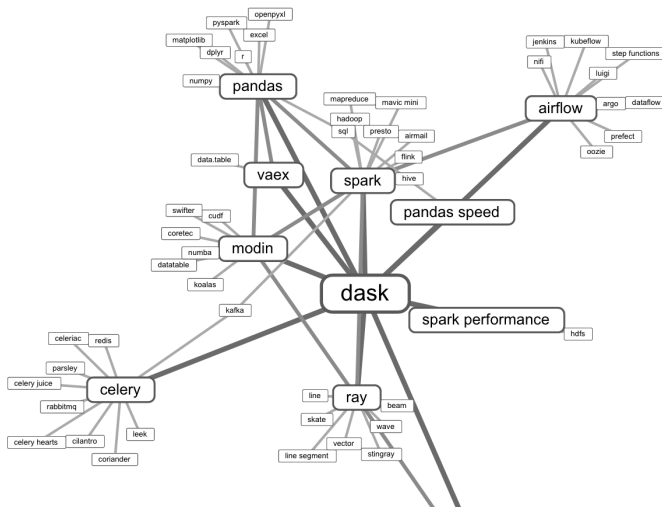
- Alternatives

- ▶ break file in multiple smaller files that can be read in parallel: useful if the reading can be done in parallel
- ▶ undersampling: need to be careful about data distribution
- ▶ use of specific hardware and software: distributed disks, distributed file systems, distributed databases, in-memory databases, parallel and distributed software
- ▶ work with compressed files: zip, parquet, CSR, CSR5 (for sparse matrices) etc

- Python libraries we will be studying:

- ▶ Dask
- ▶ Modin
- ▶ Vaex
- ▶ Koalas
- ▶ Ray
- ▶ ...

# Alternative libraries for big data



source: <https://www.datarevenue.com/en-blog/pandas-vs-dask-vs-vaex-vs-modin-vs-rapids-vs-ray>



# Differences in libraries

Function	Pandas/Modin/Koala	Vaex	Dask DataFrame	Turicreate	H2O	PySpark	Datatable
Read file	pd.read_csv() pd.read_parquet()	vaex.read_csv() vaex.open()	dd.read_csv() dd.read_parquet()	tc.read_csv() tc.SFrame()	h2o.upload_file() h2o.import_file()	sqlContext.read.csv() sqlContext.read.parquet()	dt.fread() dt.open()
Count	len(df)	len(df)	len(df)	len(df)	len(df)	df.count()	df.shape[0]
Mean	df.x.mean()	df.x.mean()	df.x.mean() .compute()	df['x'].mean() tc.Sketch(df['x']).mean()	df['x'].mean()	df.select(f.mean('x')) .collect()	df[:, dt.mean(dt.f.x)]
Standard deviation	df.x.std()	df.x.std()	df.x.std() .compute()	df['x'].std() tc.Sketch(df['x']).std()	df['x'].sd()	df.select(f.stddev('x')) .collect()	df[:, dt.sd(dt.f.x)]
Sum columns	df['x']+df['y'] df.x + df.y	df['x']+df['y'] df.x + df.y	df['x']+df['y'] df.x + df.y	df['x']+df['y']	df['x']+df['y']	df['x']+df['y']	df[:, f.x + f.y]
Sum columns mean	(df.x + df.y).mean()	(df.x + df.y).mean()	(df.x + df.y).mean() .compute()	(df['x'] + df['y']).mean()	(df['x'] + df['y']).mean()	df.select(f.mean(df['x'] + df['y'])) .collect()	df[:, dt.mean {f.x + f.y}]
Value counts	df.x.value_counts()	df.x.value_counts()	df.x.value_counts() .compute()	df['x'].value_counts()	df['x'].table()	df.select('x').distinct() .collect()	df[:, dt.count(f.x, 'x')]
Group-by	df.groupby(by='z') .agg({ 'x': ['mean', 'std'], 'y': ['mean', 'std']})	df.groupby(by='z') .agg({ 'x': ['mean', 'std'], 'y': ['mean', 'std']})	df.groupby(by='z') .agg({ 'x': ['mean', 'std'], 'y': ['mean', 'std']}) .compute()	df.groupby('z', operations={ 'c1': tc.aggregate.MEAN('x'), 'c2': tc.aggregate.STD('x'), 'c3': tc.aggregate.MEAN('y'), 'c4': tc.aggregate.STD('y')})	df.group_by('z') .mean(col = ['x', 'y']) .sd(col = ['x', 'y']) .get_frame()	df.groupby('z') .agg(f.mean('x'), f.stddev('x'), f.mean('y'), f.stddev('y'))	aggs = { 'c1': dt.mean(f.x), 'c2': dt.sd(f.x), 'c3': dt.mean(f.y), 'c4': dt.sd(f.y), df[:, aggs, dt.by(f.z)]
Join	df.join(other, on = 'key') pd.merge(df, other)	df.join(other, on = 'key')	dd.merge(df, other)	df.join(other, on = 'key')	df.merge(other)	df.join(other, on = 'key')	other.key = 'key' df[:, dt.join(other)]

source: <https://towardsdatascience.com/>

beyond-pandas-spark-dask-vaex-and-other-big-data-technologies-battling-head-to-head-a453a1f8cc13

Let's start with dataflows...

- Provides a place to run Apache Beam jobs on the GCP
- Offers the ability to create jobs based on **templates**
- No need to address common aspects of running jobs on a cluster:
  - ▶ load balancing
  - ▶ scaling number of workers for a job
- These tasks are done automatically for both **batch** and **streaming**

# Apache BEAM – Batch + strEAM

- Evolution of Google Dataflow that separates the dataflow logic from the programming issues (language, runners etc)
- Unified model for both batch and stream data processing
- Programs can be executed in different processing frameworks (via runners) using a set of different IOs

# Apache BEAM – Batch + strEAM

Why to use BEAM instead of only Hadoop, Spark, Flink, GCP Dataflow etc?

→ The Apache Beam framework provides an abstraction between your application logic and the big data ecosystem

# Apache BEAM – Batch + strEAM

In the BEAM ecosystem:

- **DataSource**: can be batches, micro-batches or streaming data
- **SDK**: Java or Python
- **Runner**: Apache Spark, Apache Flink, Google Cloud Dataflow, among others and DirectRunner (runs locally)
- To build a BEAM logic: Pipeline, PCollection, PTransform, ParDO and DoFn

# Apache BEAM – Batch + strEAM: Components

- **Pipeline:** encapsulates the workflow of your entire data processing tasks from start to finish. Includes:
  - ▶ reading input data
  - ▶ transforming that data (almost all data transformations are supported including database operations)
  - ▶ writing output data
- All Beam driver programs must create a Pipeline
- When you create the Pipeline, you must also specify the execution options that tell the Pipeline where and how to run
- Beam can run independently of the Google Cloud Platform

# Apache BEAM – Batch + strEAM: Components

- **PCollection**: distributed data set that your Beam pipeline operates on
- data may come from a fixed source like a file, or from a continuously updating source via a subscription or other mechanism



# Apache BEAM – Batch + strEAM: Components

- **PTransform**: represents a data processing operation, or a step, in your pipeline
- Every PTransform takes one or more PCollection objects as input, performs a processing function that you provide on the elements of that PCollection, and produces zero or more output PCollection objects.

# Apache BEAM – Batch + strEAM: Components

- **ParDo**: for generic parallel processing
- similar to the “Map” phase of a Map/Shuffle/Reduce-style algorithm
- a ParDo transform considers each element in the input PCollection, performs some processing function (your user code) on that element, and emits zero, one, or multiple elements to an output PCollection.

# Apache BEAM – Batch + strEAM: Components

- **DoFn**: applies your logic in each element in the input PCollection and lets you populate the elements of an output PCollection
- to be included in your pipeline, it's wrapped in a ParDo PTransform.