

Data Mining II / Adv. Topics in Data Science

Information Retrieval

Álvaro Figueira

Rita Ribeiro



1

Summary


1. [Basic Concepts](#)
2. [Text Preprocessing](#)
3. [Retrieval Models](#)
4. [Retrieval Evaluation](#)
5. [Query Modification](#)
6. [Web Retrieval](#)

2

Basic Concepts

3

Data Mining - a structured view



4

Information Retrieval and Big Data

"Big data is a term that describes large **volumes** of high **velocity**, complex and **variable** data that require advanced techniques and technologies to enable the capture, storage, distribution, management, and analysis of the information." (TechAmerica Foundation's Federal Big Data Commission, 2012)

- Big data has three dimensions described by the Three V's [Gandomi and Haider, 2015]: **Volume, Velocity and Variety**.
- **Variety** → structural heterogeneity in data
 - structured data (tables in relational databases) represent only 5% of the existing data;
 - non-structured data: text, images, audio and videos represent all the remaining data, which do not have the structure required by a computer for the analysis.

5

Information Retrieval: Definitions

- "Information retrieval **deals with** the **representation, storage, and access to documents** or representatives of documents (document surrogates)" - Gerard Salton
- "Information retrieval (IR) **is the activity of obtaining information resources relevant** to an information need from a **collection of information resources**" - Wikipedia
- "Information retrieval (IR) is **finding material (usually documents) of an unstructured nature (usually text)** that satisfies an information need from within large **collections** (usually stored on computers)" - Manning & Raghava

6

Information Retrieval

- IR helps users finding information that matches their needs.
- This process involves:
 - acquisition of information
 - organization of information
 - storage of information
 - retrieval of information
 - distribution of information
- **Goal: predict** given an information source, and apriori-knowledge from the user, **the objects which are the most relevant for the user.**

7

Information Retrieval - architecture model

General Architecture of a Information Retrieval Model

1. user poses a **query**;
2. the query is sent to the **information retrieval system (IRS)**;
3. which uses the **document index**;
4. to get **documents** with query terms;
5. compute **relevance** of documents
6. **rank** results


8

Information Retrieval (example)

The screenshot shows the ScienceDirect website interface. At the top, there's a navigation bar with 'ScienceDirect' and 'Home'. Below it, a search bar contains the text 'COVID-19'. To the right of the search bar, there are buttons for 'Search' and 'Advanced Search'. Below the search bar, there are filters for 'Additional Criteria' including 'Limit search to', 'Article type', 'Author', and 'Publication date'. The main content area displays search results for 'COVID-19' in all ScienceDirect and all Article Types, showing 170 results. The first result is titled 'Combating COVID-19—The role of robotics in managing public health and infectious diseases' by Wang, Peng, et al. (Mar 2020, 2020). The second result is titled 'Antibiotic treatment for COVID-19 complications could help resistant bacteria' by Wang, Peng, et al. (Mar 2020, 2020). The third result is titled 'COVID-19 has not caused more deaths in the US' by Wang, Peng, et al. (Mar 2020, 2020). The fourth result is titled 'COVID-19 has not caused more deaths in the US' by Wang, Peng, et al. (Mar 2020, 2020).

9

Information Retrieval (schematic process)



- How to represent **documents**?
- How to represent **queries**?
- How to **compare documents** with queries, such that the **most relevant** documents to the queries are selected?


10

Text Preprocessing

11

Text Preprocessing

To get from a documents collection a set of terms, the following text preprocessing steps are usually followed.



12


Text Preprocessing: **Lexical Analysis**

- Process by which a sequence of characters is separated in tokens.
- **Tokens** are groups of characters with a collective meaning (words).
- Word identification can be performed by means of:
 - blank spaces
 - punctuation marks: usually ignored
 - capitalization: convert to lower case
 - digits: usually ignored, but can be kept depending on the context (e.g. dates, H2O, CC4061)
 - hyphens: ignored (e.g. on-line → online) or kept (e.g., in the Knuth-Morris-Pratt algorithm)

13

Text Preprocessing: **discriminant words**

- Some words occur very frequently.
- A medium number of words present a medium frequency.
- Lots of words occur rarely.
- Those very **frequent words are not good discriminant** of the content of the documents.
- Words the **occur very rarely** are specific to particular topic, or can be **jargon, proper nouns**, etc.
- Those with **medium frequency are the most discriminant**.



14

Text Preprocessing: Stop words removal

- **Stop words**: are those with no meaning.
- Parts of speech that usually don't have meaning: pronouns, prepositions, conjunctions, determiners, ...
- Very frequent words (typically, represent more than 80% of the total number of words) and are not useful for retrieval.

- If removed, space is saved, and retrieval efficiency gained.
- The remaining is usually stored in a fast data structure (e.g., a Hash table)

15

Text Preprocessing: Stop words Removal (cont.)

- They depend on the language
 - In portuguese:
 - *a, agora, ainda, alguém, algum, alguma, algumas, alguns, antes, ao, aos, após, aquela, aquelas, aquele, aqueles, aquilo, as, até, através, cada, com, como, da, daquele, daqueles, das, de, dela, delas, dele, deles, depois,....*
 - In english:
 - *a, about, above, across, after, again, against, all, an, and, any, anybody, anything, are, as, at, be, because, been, before, being, below, between, both, but, by, ...*
- "To be or not to be"
 - A short list of stop words should be used for general collections or for inexperienced users.
 - Longer stop words lists for very specific domains.

16

Text Preprocessing: **Stemming**

- Process by which the *lexical stem* is extracted from a word.
- **Stem**: part of the word resulting when **removing affixes** (suffixes, infixes, prefixes)
- The same stem will represent a family of words semantic and morphologically related
 - {connected, connecting connection, connections, disconnected} → connect
 - {computer, computational, computation} → comput
 - {compressed, compression} → compress
- The morphologic analysis to extract the stem depends on the language and is usually complex

17

Text Preprocessing: **Stemming** (cont.)

- **Affix removal algorithms**:
 - based on heuristic methods;
 - successively, applying rules to the words;
 - different words could generate the same stem;
 - different methods: Lovins, Slaton, Dawson, Porter.
- "Porter Stemmer"
 - mainly **suffix stripping**;
 - example: plural removal using the rule with longest
 - **suffix rule example**

sses → ss	caresses → caress
ies → i	ponies → poni
ss → ss	caress → caress
s →	cats → cat

18

Text Preprocessing: **Lemmatization**

- Transformation of the word (inflectional and derivationally related forms) to the corresponding **lemma** ("the common base")
- Example of rules:
 - Convert any verbal form to infinitive:
 {am, are, is} → be
 {going, gone, went} → go
 - Plural to singular: cars → car
 - Female gender to male gender: menina → menino
- The strip will always **result in an actual word**.
- More precise than stemming but needs more resources (usually dictionaries) and more disciplines as:
 - Natural Language Processing and
 - Computational Linguistics.

19


Text Preprocessing: **Lemmatization and Stemming**

- Lemmatization vs Stemming:
 - developing a stemmer is far simpler than building a lemmatizer;
 - for lemmatizer deep linguistics knowledge is required; but, the noise will be reduced, and the information retrieval process will be more accurate.
- Both methods are used to reduce the size of the vocabulary.
- Not all the words are indexed but only a kind of representatives.
- Advantages:
 - reduction in the vocabulary → efficiency and space saving;
 - the number of matching documents to retrieve increases.
- Disadvantages:
 - information about the complete word is lost.

20


Text Preprocessing

Tokenization → Stop words removal → Lemmatization



21

Convered so far...



Remember: this can be done on a set of documents


22

Indexation of terms

Objective: To build an **Index** from a set of documents

Inverted index construction

Index = Dictionary + Posting list



- efficient data structure (e.g. Hash Table, B-Tree, Trie)
- fast access to the collection to get the relevant documents for a query
- terms are stored, pointing to the documents where they occur, and their corresponding weight.

Retrieval Models

Information Retrieval Models

- Each document is seen as the set of its terms.
- A **term** may represent a single word or multiword units (e.g., "White House").
- **Weights can be associated to terms.**
- Different models find **weights** in different ways
 - Boolean Model
 - Vector Space Model
 - Binary Scheme
 - TF Scheme
 - TF-IDF Scheme
 - Statistical Language Model (uses Bayesian reasoning)

} We will study these

21

25

Information Retrieval Models: Boolean Model

- The simplest retrieval model.
- Each document is represented by the presence/absence of terms.
- A **query** is a boolean expression with **and**, **or** and **not** operators to join terms.
- It looks for **exact match**, i.e. every document that matches the query.
- **Result**: set of documents, without ordering, satisfying the query.

document	t_1	t_2	t_3
d_1	1	0	1
d_2	1	0	0
d_3	0	1	1
d_4	1	1	1

query:

$$t_1 \wedge (t_2 \vee \neg t_3)$$

result: { d_2, d_4 }

26

Information Retrieval Models: Boolean Model (cont.)

- **Advantages:**
 - very efficient;
 - predictable, easy to explain;
 - it works well with experienced users;
 - it is still used in easy searches (e.g., email).
- **Disadvantages:**
 - returned documents are not represented by relevance;
 - "close" documents are not presented as such;
 - it is difficult to formulate complex boolean queries for normal users;
 - they can return too many documents, or very few.

27

Information Retrieval Models: **Vector Space Model**

- **Terms** are **axes** of the space.
- Documents and queries are **vectors** in this space.
- Vector coordinates are term **weights**.
- Very high-dimensionality but very sparse vectors (most entries are zero)
- Ranks the documents by similarity degree (e.g., cosine)
- Allows to retrieve documents that partially match the query.



28

Information Retrieval Models: Vector Space Model (Binary)

Binary Scheme

- The weights indicate the **presence or the absence** of a term in a document.

$$w_{bi}(t_j, d_i) = \begin{cases} 1 & \text{if } t_j \in d_i \\ 0 & \text{otherwise} \end{cases}$$

29

Information Retrieval Models: **Term-Frequency****Term-Frequency (TF) Scheme**

- Motivation:** A term appearing often in a document may be more important for identification than a term appearing rarely.
- This aspect is captured by **tf**, the "term frequency" in a document.

$$w_{TF}(t_j, d_i) = tf(t_j, d_i)$$

where **tf** (t_j, d_i) is the **frequency** of term j in document i

- It can be normalized:

$$w_{TF}(t_j, d_i) = \frac{tf(t_j, d_i)}{\sum_{t_k \in d_i} tf(t_k, d_i)}$$

30

Information Retrieval Models: Term-Frequency (cont.)

Term-Frequency (TF) Scheme

Example:

document	t_1	t_2	t_3	$w_T(t_1, d)$	$w_T(t_2, d)$	$w_T(t_3, d)$
d_1	1	0	1	1/2	0	1/2
d_2	1	0	0	1	0	0
d_3	1	1	0	1/2	1/2	0
d_4	2	1	1	1/2	1/4	1/4
	5	2	2			

31

Information Retrieval Models: Inverse Document Frequency

Inverse Document Frequency (IDF)

- **Motivation:** If a term appears many times, in many documents, it will probably be irrelevant. However, if a term occurs in just a few documents it will have a greater discriminating power.
- This aspect is captured by *idf*, the "inverse document frequency"

$$idf(t) = \log \left(\frac{N}{|\{d \in D: t \in d\}|} \right)$$

Where, N is the number of documents in the corpus (D).Therefore, we are dividing it by the number of documents in which term t occurred.

32

Information Retrieval Models: **TF-IDF****Term Frequency - Inverse Document Frequency (TF-IDF) Scheme**

- The weights take into account both intra-document and inter-document term frequency.

$$w_{TF-IDF}(t_j, d_i) = tf(t_j, d_i) \times idf(t_j) = \frac{tf(t_j, d_i)}{\sum_{t_k \in d_i} tf(t_k, d_i)} \times \log \left(\frac{N}{|\{d \in D: t_j \in d\}|} \right)$$

Example:

document	t_1	t_2	t_3	TF_{t_1}	TF_{t_2}	TF_{t_3}	$W_{TF-IDF}(t_1, d_i)$	$W_{TF-IDF}(t_2, d_i)$	$W_{TF-IDF}(t_3, d_i)$
d_1	1	0	1	1/2	0/2	1/2	0	0	$1/2 \times \log(2)$
d_2	1	0	0	1	0/1	0/1	0	0	0
d_3	1	1	0	1/2	1/2	0/3	0	$1/2 \times \log(2)$	0
d_4	1	1	1	1/3	1/3	1/3	0	$1/3 \times \log(2)$	$1/3 \times \log(2)$
$df(t_i)$	4	2	2						
$N/df(t_i)$	1	2	2						
$\log(N/df(t_i))$	0	$\log(2)$	$\log(2)$						

33

Information Retrieval Models: **Document Similarity****Similarity for retrieval**

- A *query* is represented in the same manner as other documents.
- Relevant documents are the ones closer to the query, using some similarity metric, e.g., [cosine similarity](#):

$$\cos(d_i, q) = \frac{\sum_{t_k \in T} w(t_k, d_i) \times w(t_k, q)}{\sqrt{\sum_{t_k \in T} w(t_k, d_i)^2} \times \sqrt{\sum_{t_k \in T} w(t_k, q)^2}}$$

Example:

document	t_1	t_2	t_3	
d_1	1	0	1	$\cos(d_1, q) = 0.5000000$
d_2	1	0	0	$\cos(d_2, q) = 0.0000000$
d_3	1	1	0	$\cos(d_3, q) = 0.5000000$
d_4	1	1	1	$\cos(d_4, q) = 0.8164966$
q	0	1	1	

34

Retrieval Evaluation

35

Retrieval Evaluation

The **quality** of an Information Retrieval System depends on:

- space for indexing the documents
- time, i.e., efficiency in retrieving and indexing
- user satisfaction in terms of usability or operation
- retrieval **effectiveness**

36

Evaluation Measures

- How to evaluate the **retrieval effectiveness**?

$$\text{precision} = \frac{\# \text{ relevant docs retrieved}}{\# \text{ docs retrieved}}$$

$$\text{recall} = \frac{\# \text{ relevant docs retrieved}}{\# \text{ relevant docs}}$$

- Objective:** get as much relevant documents as possible while at the same time, getting as few irrelevant documents as possible

37

Evaluation Measures (summary)

Summary of the most used measures

- prec@k:** precision at top k retrieved documents
- rec@k:** recall at top k retrieved documents
- MAP** (Mean Average Precision): average precision each time a relevant document is retrieved
- F1:** harmonic mean between precision and recall: $F1 = 2 \times \frac{P \times R}{P + R}$
- precision-recall curves**
- NDCG** (Normalized Discounted Cumulative Gain): gain, or graded relevance/usefulness of a document, is accumulated starting at the top of the ranking and may be reduced, or discounted, at lower ranks

38

The difference between micro, macro, and weighted averages

- In multi-class classification problems, models often compute a metric for each class.
- Example: in a 3-class problem, three precision scores are returned.
- However, we just need a single global metric – the averaging methods

1. Macro average

A **simple arithmetic mean**. Example: if precision scores are 0.7, 0.8, 0.9, macro average would be their mean = 0.8.

2. Weighted average

This method **takes into account the class imbalance** as metrics for each class are multiplied by the proportion of that class. Example: if there are **100 samples** (30, 45, 25 for each class respectively) and the precision scores are .7, .8, .9, the weighted average would be:

$$0.3 * 0.7 + 0.45 * 0.8 + 0.25 * 0.9 = 0.795$$

3. Micro average

Micro average is the same as accuracy – it is calculated by **dividing the number of all correctly classified samples** (True Positives) **by the total number of correctly and incorrectly classified** (True Positives + False Positives) samples of each class.

Obs: Use micro average when there is an imbalanced problem. This approach does not take into account class distribution/contributions. Which means it is sensitive to performance on rare classes.

Evaluation Measures (cont.)

Example: the system ranked the documents like this...

Rank	Relevant
1	+
2	+
3	-
4	+
5	-
6	-
7	+
8	-
9	-

(#Relevants = 4)

$$prec@2 = 1 \quad prec@4 = 3/4 \quad prec@8 = 1/2$$

$$rec@2 = 1/2 \quad rec@4 = 3/4 \quad rec@8 = 1$$

$$MAP = (1/1 + 2/2 + 3/4 + 4/7)/4 = 0.8303571$$

$$F1@2 = 2 * (1 * 1/2)/(1 + 1/2) = 0.6666667$$