

Index

- [IPv6](#)
- [ICMPv6](#)
- [Other Layers](#)
- [Transition and Interoperation Mechanisms](#)

IPV6

Why IPv6?

- Address Exhaustion: IPv4 addresses are running out; NAT is a temporary solution
- Improve router performance: simplifies IP header, aligns to 64 bits, address hierarchy with more levels, and simplifies routing tables
- Improve Mobile IP support

Headers

- **Ver** - protocol version (=6)
- **Class** - similar to IPv4's ToS / DSCP
- **Flow Label** - flow identification;
 - Unique for each flow with the same (src,dst) pair
 - Zero means "no flow label"
- **Payload length** - data field size
 - Limited to $2^{16}-1$ bytes, but there there's an option for Jumbograms up to $2^{32}-1$ bytes
 - No need for total length since the header size is fixed
- **Next Header** - type of next header
- **Hop Limit** - IPv4's TTL (but done properly)
- No checksum, with error checking performed by lower layers

| | | | | |
|---------------|---------------------------|-----------|-----------------|---------------|
| V6 (40 bytes) | Ver(4) | Class (8) | Flow Label (20) | |
| | Payload Length (16) | | Next Header (8) | Hop Limit (8) |
| | Source Address (128) | | | |
| | Destination Address (128) | | | |

| | | | | | |
|----------------|--------------------------|--------|--------------|----------------------|------------------|
| V4 (20+ bytes) | Ver(4) | IHL(4) | DSCP (8) | Total Length (16) | |
| | Identification (16) | | | Flags(4) | Frag Offset (12) |
| | TTL (8) | | Protocol (8) | Header Checksum (16) | |
| | Source Address (32) | | | | |
| | Destination Address (32) | | | | |
| | Options | | | | Padding |
| | | | | | |

Addresses

- 128 bits represented in hexadecimal
- Leading zeroes may be omitted
- A sequence of zeroes may be abbreviated by ::
 - Only one :: in an address to avoid ambiguity
- Mask (prefix length) similar to IPv4 CIDR
- IPv4 can be embedded in IPv6 addresses


Unicast

- Loopback ::1 (/128)

- Link-local:
 - Only used on the local link (not routed)
 - FE80::/64 (last 64 bits are device identifier)
- Site-local (**deprecated**):
 - Used within the same site, and therefore only valid therein
 - FEC0:0000:0000:<Subnet(16)>:<Interface(64)>
- Global address:
 - Internet routable address
 - 2000::/3
 - But there can be routable IPv6 addresses outside this range

Multicast

- Multicast:
 - With different scopes



- 4-bit multicast scope field used to limit the scope of the multicast group

| | | |
|-------------------------|----------------------------|----------------|
| 0 Reserved | 6 (unassigned) | C (unassigned) |
| 1 Interface-Local scope | 7 (unassigned) | D (unassigned) |
| 2 Link-Local scope | 8 Organization-Local scope | E Global scope |
| 3 Realm-local scope | 9 (unassigned) | F Reserved |
| 4 Admin-Local scope | A (unassigned) | |
| 5 Site-Local scope | B (unassigned) | |

Other

- Anycast address
 - Defined on more than one interface, but delivered only to one
- Unspecified address
 - Denoted by :: and used only during configuration

Required addresses

- Node:
 - Loopback
 - Link-local for each interface
 - Configured unicast (or anycast)
 - All-nodes multicast
 - Solicited-node multicast address for each of its unicast and anycast addresses
 - Multicast groups that the node belongs to
- Router: all the above plus
 - Subnet-Router Anycast (equal to the subnet prefix)
 - All-Routers multicast

ICMPv6 (Internet control message protocol for IPv6)

-
- **Type** defines the type, **code** a subtype (similar to ICMP (v4))
 - Example:
 - Type = 1 (Error code for destination unreachable)
 - Code = 0 (no route to destination)
 - Errors have type with high-order bit = 0 (0-126)
 - Informational are 128-254

Neighbour Discovery

- Replaces IPv4 ARP
- Adds auto-configuration helper functions
- Extensible: messages may have options for added information
- IP Hop limit: set to 255

Neighbour solicitation

- Sent by hosts or routers, for address resolution, neighbour unreachability detection (NUD), and duplicate address detection (DAD)
- Type = 135, code = 0
- **Target address**
 - IP address of the solicitation target
- **IP destination address**
 - solicited-node multicast address corresponding to the target address, or the target address itself
- **IP source address**
 - address of the interface where the packet is sent, or unspecified address for DAD

Neighbour advertisement

- Sent by hosts or routers, in response to solicitations or unsolicited
- Type = 136, code = 0
- **Target address**
 - For solicited: same as in solicitation
 - For unsolicited: IP address that changed link-layer address
- Option: source link-layer address of sender

Solicited-Node MC Address

- Multicast address formed using ff02::0:0:0:1:ff00::/104 prefix and low-order 24 bits of an address (unicast or anycast)
 - Example:
 - Unicast address: 4037::01:800:200e:8c6c
 - Solicited Address: ff02::0:0:0:1:ff0e:8c6c
- Most probably, no two nodes in the same network map to the same multicast address → much more efficient than the local broadcast used by ARP!

Router Discovery

- For discovering routers on the link and, possibly, address information

- **Router solicitations** (Type = 133, code = 0)
 - Sent by hosts to the link-local all-routers multicast address
- **Router advertisements** (Type = 134, code = 0)
 - Sent by routers
 - In response to solicitations - to unicast address of requester
 - Unsolicited - to all-nodes multicast address with link-local scope
 - Advertise: net prefix, default router(s), MTU

Auto Configuration - SLAAC

- **StateLess** Address AutoConfiguration
 1. Generating link-local address
 2. Generating global address(es)
 3. Performing duplicate address detection (DAD)
- Based on Router Advertisements
- Uses advertised prefix and MAC address of interface
- Valid only for /64 networks

What about DNS servers?

- IPv6 Router Advertisement Options for DNS Configuration
- Type: 25
- Length: indicates the length of the option
 - Implicitly the number of advertised DNS servers

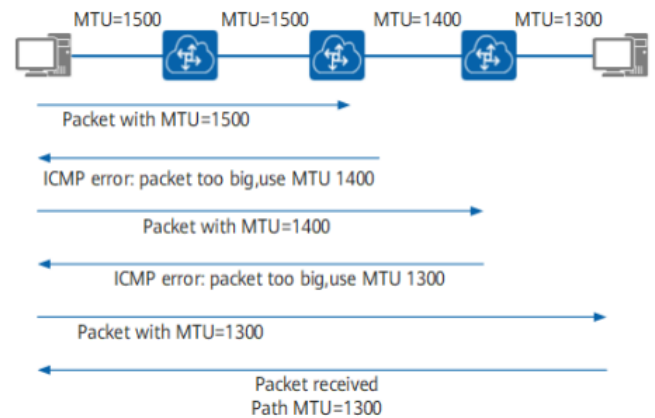
Duplicate Address Detection

- Used after address auto-configuration to check uniqueness of address
 - Not performed for anycast addresses (configured)
- Node sends neighbour solicitation with:
 - **Src address:** unspecified
 - **Dst address:** solicited-node multicast address of the tentative address
 - **Target address:** the tentative address
- Receiving Neighbour Solicitation
 - Target address is the tentative address
 - If src address is unicast (not DAD): ignore packet
 - If src address is unspecified: other node wants to use same address (DAD)
 - No node will use the target address (as extra precaution)
- Receiving Neighbour Advertisement
 - Target address is the tentative address
 - Address is not unique (cannot be used)

Path MTU Discovery

- IPv6 does not support fragmentation in routers
- Senders need to transmit packets \leq smallest MTU along the path (i.e., Path MTU)
- Procedure for discovering the Path MTU
 - Source node assumes path MTU = 1st hop link MTU

- If this is too large for a link
 - router drops the packet
 - returns an ICMPv6 Packet Too Big (type = 2, code = 0)
- Source reduces MTU to the indicated size
- When the MTU used by the source \leq Path MTU, packets reach the destination



Other Layers

DHCPv6

- Stateful configuration
- Uses UDP (like DHCPv4)
- DHCP clients use the link-local address
- Uses two multicast groups
- Requires DAD to ensure address uniqueness
- Possible to obtain partial information via Information-Request message
- DHCP Unique Identifier (DUID) for client-server identification
 - Not all messages require this ID
- Optional authentication and encryption
- Information carried in options

Messages

| DHCPv6 Message Type | DHCPv4 Message Type |
|------------------------------------|---------------------|
| Solicit (1) | DHCPDISCOVER |
| Advertise (2) | DHCPOFFER |
| Request (3), Renew (5), Rebind (6) | DHCPREQUEST |
| Reply (7) | DHCPACK / DHCPNAK |
| Release (8) | DHCPRELEASE |
| Information-Request (11) | DHCPINFORM |
| Decline (9) | DHCPDECLINE |
| Confirm (4) | none |
| Reconfigure (10) | DHCPFORCERENEW |
| Relay-Forw (12), Relay-Reply (13) | none |

Recursive DNS server configuration

- Approaches for DNS configuration in clients
 - Router Advertisements, DHCPv6, or well-known anycast address

DNS - changes for IPv6

- AAAA new record type
 - IPv6 address
- AAAA queries returns all IPv6 addresses associated with domain name

- IP6.ARPA domain for reverse (PTR) queries (similar to IPv4's in-addr.arpa)
- Existing query types
 - NS, SRV, MX: re-defined to also return AAAA entries

DNS in IPv4 and IPv6

- Maintain IPv4 and IPv6 accessible DNS recursive server
 - IPv4 only or dual stack
 - At least one IPv4 reachable server per DNS zone

Transition and Interoperation Mechanisms

IPv6 in an IPv4 world

Problem

- Incompatibility arises due to differences in header formats, Sockets API, and applications


Solutions

Dual stack nodes

- Support both IPv4 and IPv6
- Routers must maintain both routing tables, protocols, etc.

Configured Tunneling (Aka 6in4)

- Tunnel endpoints encapsulate the original IPv6 in an IPv4 packet
- Tunnel endpoints (routes) are manually configured
- Does not scale well
 - Requires a different tunnel for each pair of routers $\rightarrow O(n^2)$




6to4

- 6to4 Tunnel (without explicit setup)
- Use of 2002::V4ADDR::/48 networks
- Sites use 6to4 relay to communicate through the tunnels
 - These relay routers can have an IPv4 anycast address of 192.88.99.1
 - Sites may also have connection to IPv6-only networks
- **Components:**

Router B**6to4 Router C****6to4 relay Router A**

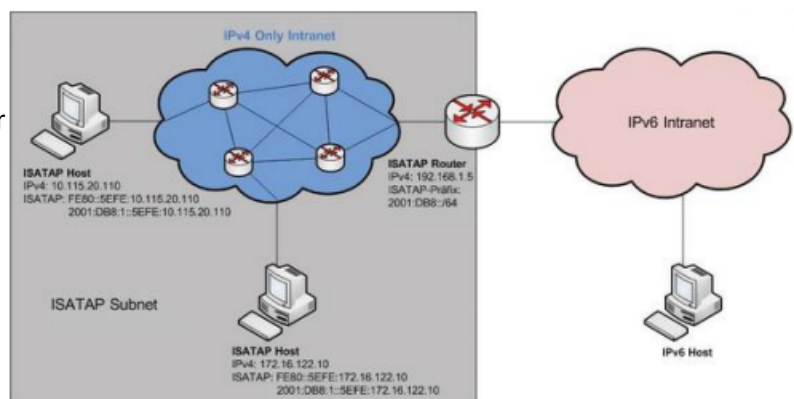
- Connects to the IPv6 backbone
 - Has a native IPv6 address (not 2002::/16)
 - Does not handle IPv4
- Connects to the IPv4 backbone
 - Does not relay packets to the IPv6 backbone
 - 6to4 IPv6 address and a native IPv4 address
- Connects to the IPv4 and IPv6 backbone
 - Has a 6to4 and a native (not 2002::/16) IPv6 address
 - Has a native IPv4 address
- **Issues:** 6to4 causes asymmetric routing and relays may receive traffic from anywhere. 6rd is a slight modification that solves this issue

**6rd**

- **IPv6 Rapid Deployment** on IPv4 Infrastructures
- Similar to 6to4 but using ISP-specific prefix instead of 2002::/16
 - Allows use of private IPv4 addresses
 - ISP-specific prefix length not fixed
- The ISP
 - Operates one or more gateways at the IPv4/IPv6 border
 - Deploys relays with specific anycast addresses (only for its customers)
 - Has more control over traffic flowing through its network

Intra-Site Automatic Tunnel Addressing Protocol (ISATAP)

- Connects dual-stack nodes over IPv4 networks
- Uses entire IPv4 network as a link layer (NBMA)
- Defines locators (mappings) on interfaces for routing packets
- Requires a Potential Router List (PRL) (obtained by doing a DNS lookup for



isatap.<domain>) for communication with the public IPv6 Internet

- ISATAP IPv6 addresses:
 - <prefix>:0000:5efe:<IPv4 addr>

Teredo

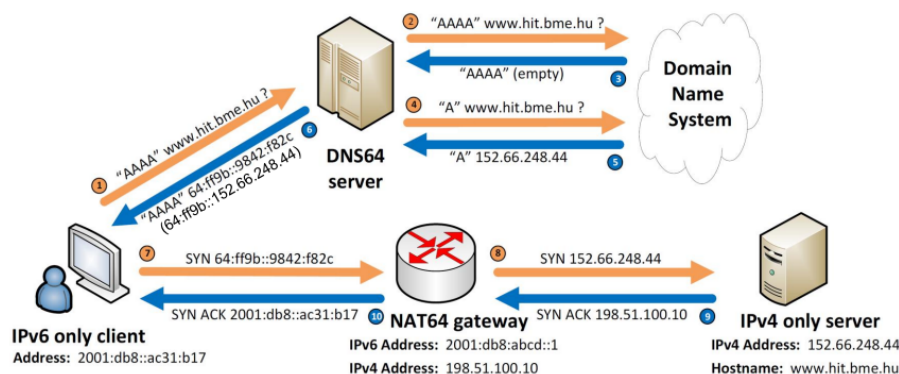
- Allows automatic IPv6 tunneling between hosts that are located across one or more IPv4 NATs
- Encapsulates IPv6 packets in UDP
 - More overhead than other techniques → use only when they are infeasible
- Use of 2001:0000::/32 for Teredo clients
 - 2001:0000:<teredo server IP address>:<flags (16)>:<obscured UDP port>:<obscured public IPv4 address>
- Discovers and maintains NAT mappings to the client
- Components:
 - Client: has IPv4 connectivity, wants IPv6
 - Server: to discover external address and type of NAT
 - Relay: forwards traffic using the Teredo encapsulation on IPv4 and to the IPv6

IPv4 in an IPv6 world

- Increasing deployment of IPv6 → need for
 - Nodes in IPv6-only networks to reach IPv4-only services on the Internet
 - Traversal of IPv6-only operator networks to reach IPv4-only services on the Internet
- Larger IPv6 addresses make it easier to use translation/NAT instead of tunneling
 - IPv6 addresses can embed IPv4 addresses

NAT64/DNS64

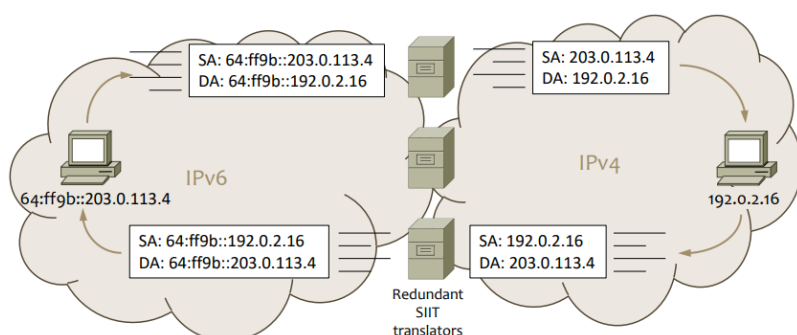
- NAT64 facilitates access to IPv4-only services from IPv6-only networks by translating IPv6 to IPv4 addresses.
- DNS64 synthesizes AAAA records for services offering only A records, using the chosen prefix.
- Example:



- Limitations include IPv4-only software, IPv4 literals (no DNS lookup is performed), and IPv4 networks behind IPv6-only ISP networks

Stateless IP/ICMP Translation (SIIT)

- Translates headers of IP packets and IP packet fragments inside ICMP messages (for transparency)
- Comes in two flavors:
 - "Traditional" SIIT (entire IPv4 address embedded in an IPv6 address with given prefix)
 - SIIT with Explicit Address Mappings (EAM) (uses configured host-specific or block mappings)
- Advantages of stateless translation:
 - No need to maintain per-flow state
 - Easier load distribution - translator for outgoing and incoming packets needs not be the same
- Disadvantage: the 1:1 mapping between IPv4 and IPv6 addresses wastes scarce IPv4 addresses
- Example ("traditional" SIIT):



464XLAT

- 464XLAT is an architecture combining Stateful NAT64 with an additional, stateless translator (SIIT)
 - CLAT: Client-side transLATOR (SIIT)
 - PLAT: Provider-side transLATOR (Stateful NAT64)
- Addresses the use cases not covered by NAT64/DNS64

